

Sovereign Anchor Holdings – The July 2026 Dialogue

Hello, everyone. Welcome to the July dialogue of Sovereign Anchor Holdings.

Before anything else, I want to flag that this one is going to sound different from what you're used to. Not just longer – different in structure. Up until now, if you were following along, you'd get two separate documents from us each quarter: a Dialogue, which was me, talking, reflecting, giving you the portfolio and the philosophy behind it – and separately, a Research Notes document, which was dense, technical, written in that flat analyst voice, meant to be read with a highlighter in hand. Two different products, two different moods.

Starting this quarter, I'm merging them. One conversation. My portfolio update, my worldview, and the research – all of it – sitting inside the same room, because I think that's actually more honest about how this works. I don't think about my portfolio for an hour and then switch hats and think about Switzerland's debt brake for a separate hour in some detached, clinical voice. It's all one continuous stream of thinking, and I'd rather hand you the stream than hand you two edited transcripts of it. So that's the experiment this quarter – merge the research into the dialogue, make it feel like one person talking you through what's actually going on in his head, rather than a report you're supposed to study. If it works, we keep doing it this way.

The Portfolio, As It Stands Today

Let's start where I always start – with the actual money, the actual positions, because everything else I say today is downstream of this.

A lot has changed since the last time we spoke. I've gone from a wider basket down to eight stocks. That's a real, deliberate move – I've gotten more concentrated, and I've also moved further down the market-cap pyramid than before, out of pure mid-cap comfort and into some genuine small-cap territory. Both of those are choices, not accidents, and I'll explain why in a minute.

The eight names, if you want the full list: Aptus Value Housing Finance, Bajaj Finance, Caplin Point Laboratories, Cholamandalam Investment and Finance, KPIT Technologies, Muthoot Finance, Tata Consultancy Services, and Waaree Energies.

By sector, it breaks down like this – financials are 51.6% of the book, software is 20.58%, healthcare is 15%, and engineering is

12.8%. By market-cap band, using the standard Nifty classification, large-cap is 50% and small-cap is 37%, with the balance in mid-cap.

And by individual weight: Cholamandalam sits at 14.33%, Bajaj Finance at 12.68%, TCS at 12.17%, Muthoot at 11.06%, Waaree at 12.81%, Caplin Point at 15%, Aptus at 13.5%, and KPIT at 8.41%.

Now, why financials at over half the book? Because that's where I think the cleanest, most durable compounding in India actually lives right now – lenders and housing financiers who are riding the country's credit-penetration story without taking on the operational complexity of building something from scratch. Why the concentration? Because when you only hold eight names, every single one of them has to earn its place, and the churn from here on should be a lot lower than it's been. I say "hopefully" very deliberately there – I'm not going to pretend I can guarantee that. Markets have a way of forcing your hand.

The core view underneath all of it hasn't moved: I want high growth, high quality, at a reasonable valuation, with a level of risk I can actually live with. And because the book is concentrated now, that risk question matters more than it used to. Eight positions means each one is a real bet, not a rounding error.

So why the shift in the first place – why go from a broader, more diversified starting point to this tighter, more concentrated one? Honestly, it comes down to risk appetite changing, and risk appetite changing for very ordinary reasons. Early on, the money itself was constrained – there just wasn't a lot of room to be adventurous. Over time, two things loosened that constraint. First, the time horizon available to me got longer, which by itself lets you underwrite more risk, because time is the thing that turns volatility into a rounding error. Second – and this one's less about the calendar and more about me – my own understanding of risk got better. The modeling, the frameworks, the pattern recognition, all of it matured. When your risk appetite actually is wider, and you can defend why, it would almost be dishonest to keep running a portfolio built for a narrower one.

Who's Actually Talking to You

I realize I've never properly introduced myself on one of these, so let me do that now, because I think it matters for how you should weigh everything I'm about to say.

My name is Anmol Verma. I'm 22. I started investing a little over three years ago. I'm a recent graduate from IIT, and alongside running this, I work inside a family office — which, if you think about it, is a strange and slightly funny arrangement, because it means my own capital is compounding in the same room where I earn my salary. My money is working for me while I'm working. I don't see a conflict there — I see an alignment. Why wouldn't I want both engines pointed the same direction?

I'm currently a CFA Level 3 candidate, though I want to be honest with you — I don't think that credential is the important part of this story. It's a useful discipline, a good filter for rigor, but it's not why you should listen to me. If anything, the more important thing to understand about how I see myself is this: I don't think of myself as an "investor" in the grand, mythologized sense. I think of myself as a minor business partner — small "m," small "p" — in eight companies I happen to have bought small slices of. That framing matters more than it sounds like it should. An investor watches a stock price. A partner asks whether the business is still doing what it said it would do. I'm trying to be the second kind of person, and I'd rather you hold me to that standard than to the first.

And because I see it that way, you shouldn't expect wild, frequent changes from me. Partners don't rotate out of businesses every quarter because the mood shifted. I'm sitting on the shoulders of people who've actually made money in this business over decades, and the humility that comes with knowing that — knowing I'm 22 and three years into this — is, I think, the single healthiest thing I can bring to how I manage your trust and mine.

Why I'm Not in a Hurry

There's a belief that sits underneath everything I do, and I'd rather say it out loud than leave you to infer it: I don't think life hands you very many genuinely important decisions. Three or four truly good calls, made with real conviction and held onto, can be life-changing. Everything else is noise dressed up as activity.

That belief is exactly why you're going to see a slower cadence from me, not a faster one, even as the portfolio gets more concentrated. Concentration sounds like it should mean more action, more monitoring, more fiddling — but for me it means the opposite. It means I've already decided these eight are the ones that matter, and now my job is to leave them alone and let compounding do the unglamorous, invisible work it does. I'm not trying to manufacture a reason to trade every quarter just because

a quarterly letter demands a narrative. If the honest update is "nothing changed, and that's good," I'll tell you that.

This is also, I think, the first place where I should tell you what I've actually been reading, because none of this – the concentration, the patience, the refusal to treat every data point as an emergency – is something I invented on my own. It's borrowed. Almost all of it is borrowed from people who thought about risk and decision-making more rigorously than I ever will on my own, and I want to walk you through that library properly, because I think understanding *how* I think will tell you more about what I'll do with your money than any single stock pick will.

The Library – What I've Actually Been Reading, and Why It Matters More Than Any Single Stock Tip

I keep a running library of the books that have actually changed how I think – not books I skimmed, books I've gone back to more than once. I want to walk you through it properly this quarter, because I think if you understand the shelf, you'll understand the portfolio decisions without me having to justify each one individually.

Let's start with the man who probably shaped my risk instincts more than anyone else – Nassim Nicholas Taleb. I've read the trilogy – *Fooled by Randomness*, *The Black Swan*, and *Antifragile* – and if I had to compress all three into one sentence, it would be this: in any domain where randomness dominates, we systematically mistake luck for skill, and then we build comforting stories to explain it afterward. Markets are exactly that kind of domain. Short-term outcomes are noisy, extreme events dominate the long-run results, and survivors end up looking like geniuses purely because we never see the graveyard of people who took the same bet and lost. Taleb calls that the silent evidence problem – we study the trader who made a hundred million and never study the thousand who blew up trying the identical strategy. That idea alone is why I refuse to judge myself, or any position in this portfolio, purely on a good quarter or a bad one. A good process can produce a bad result. A bad process can produce a good result. Outcome is not the same thing as decision quality, and if you only remember one sentence from this entire library, let it be that one.

The other Taleb idea that sits underneath this whole portfolio is the distinction between fragile, robust, and antifragile – things that break under stress, things that survive stress, and things that actually get stronger from it. It's why I care so much about balance sheets with almost no debt, why a zero-leverage compounder appeals to me more than a leveraged one growing faster on borrowed time. Debt is exactly the kind of thing that looks fine in

Mediocristan — a world of averages — and then kills you in Extremistan, a world where one tail event can erase a decade of gains. I'd rather own eight boring, antifragile businesses than twenty exciting, fragile ones.

Then there's Howard Marks — *The Most Important Thing* and *Mastering the Market Cycle*. Marks is where I got the phrase "second-level thinking," and it's become a genuine filter I run every idea through. First-level thinking says "this is a good company, buy it." Second-level thinking asks what's already priced in, what the market currently believes, and where that consensus might be wrong. You cannot beat the market by thinking the same thing everyone else is thinking, even if you're thinking it correctly — being right and being non-consensus at the same time is the actual job. Marks is also where the cycle-awareness comes from. Nothing goes in a straight line — not credit, not sentiment, not commodity prices, not even good businesses' valuations — and pretending otherwise is how people get run over holding the right stock at the wrong price.

Ray Dalio gave me the macro lens I use to think about entire countries, which matters enormously this quarter because, as I'll get to, we spent a lot of research hours this quarter on country-level work. *Principles for Dealing with the Changing World Order* and *How Countries Go Broke* both build around the idea of a "Big Cycle" — a fifty-to-hundred-year arc that empires move through: rise, peak, decline, and then conflict and restructuring. Dalio tracks eight measurable determinants of national power — things like education, innovation, competitiveness, share of world trade, financial center strength, and reserve currency status — and the pattern he shows again and again is that these things rise together and decline together. A country doesn't just wake up one day less competitive. It's usually been visible in the data for a decade before the narrative catches up. That's exactly the lens I want on before I ever look at a country's stock market — is this a nation in the early stage of its internal cycle, still doing the unglamorous work of sound finances and rising productivity, or is it further along, living off accumulated prosperity while the foundations quietly erode?

Ruchir Sharma's *Rise and Fall of Nations and Democracy on the Road* sit right next to Dalio on my shelf, and honestly do a lot of the same job from a different angle. Sharma's whole argument is that growth is never linear — countries don't rise or fall because of some fixed destiny, they rise and fall because reform momentum comes and goes. A crisis or a change in leadership triggers reform, growth accelerates, and then — this is the part people forget — success breeds complacency. Elites relax, structural reform slows, corruption creeps back in, and the growth that felt

permanent turns out to have had an expiration date the whole time. He's also sharp on debt as what he calls a "growth mirage" – credit expansion can make an economy look prosperous right up until it doesn't, and rapid private credit growth is one of the most reliable predictors of a coming slowdown that conventional forecasters miss because they're staring at GDP and demographics instead of the direction of change.

For the actual mechanics of thinking clearly, I keep coming back to Peter Bevelin's *Seeking Wisdom: From Darwin to Munger*. It's not motivational, it's structural – a manual on how to avoid being stupid in complex systems, built by borrowing from biology, physics, psychology, and Charlie Munger's mental models all at once. The core idea is that reality is probabilistic, not certain, and most of our errors come from ignoring randomness, confusing correlation with causation, or oversimplifying systems that are genuinely nonlinear. The Darwinian piece stuck with me particularly hard – survival isn't about being the strongest or the smartest, it's about fitness to the specific environment, and businesses are exactly the same. A competitive advantage that doesn't evolve is a competitive advantage on a countdown timer. That's part of why, when I look at a company like Caplin Point or Dixon Technologies, I'm not just asking "is this a good business today," I'm asking whether the moat is the kind that adapts or the kind that's already fossilizing.

On the more civilizational end of the shelf, I've got Yuval Noah Harari's *Sapiens* and *Nexus*, and Jared Diamond's *Guns, Germs, and Steel*. These aren't investing books at all, but I think they're some of the most important books on the shelf precisely because they're not. They force you to zoom out to a timescale where "quarterly earnings" doesn't even register – where the actual question is why some societies accumulated the institutions, the shared stories, the capacity for large-scale cooperation, that let them compound advantage over centuries while others didn't. When I'm evaluating a country as a place to deploy capital – and this quarter that meant Switzerland and Australia specifically – I'm partly asking a Harari-and-Diamond question underneath the Dalio-and-Sharma question: does this place have the deep institutional and geographic inheritance that makes stability the default, or is the current stability more recent and more fragile than it looks?

Adam Smith's *The Wealth of Nations* is the oldest book on the shelf, and I still think it's underrated as an investing text, mostly because everyone assumes they already know what it says. Division of labor, specialization, the invisible hand – these ideas get reduced to slogans, but reading the actual argument is a good discipline in understanding why some economies specialize into

genuinely difficult-to-replicate niches – precision engineering, specialized financial services – while others stay commoditized. It's basically the 18th-century ancestor of everything Dalio and Sharma are measuring today. And Adrian Wooldridge's *What Went Wrong with Capitalism* is a good modern counterweight – a reminder that the system Smith described isn't static, it evolves, and sometimes it evolves in directions that create real structural problems worth taking seriously rather than dismissing.

On the numbers side, nobody has shaped how I actually build a thesis more than Aswath Damodaran – *Narrative and Numbers* and *Corporate Life Cycle*. His whole point is that great investing sits at the intersection of a good story and good numbers, and that people fail in one of two predictable ways: story-driven investors who dismiss valuation models as backward-looking mechanical nonsense, or spreadsheet-driven investors who build detailed models around narratives that don't actually hold up. A real thesis has to answer what the company's story is, how that story turns into actual revenue, margin, and reinvestment numbers, and whether those implied numbers are even plausible. Narrative sets the assumptions. Numbers test whether the narrative is credible or just a nice story. Every single one of the companies I looked at this quarter – and there were a lot of them – got run through exactly that filter: what's the story, and do the numbers actually support it, or are we just in love with the story?

Peter Bernstein's *Against the Gods* is the history-of-risk book I'd give to anyone who thinks probability is boring. His core claim is that the defining achievement of modern civilization isn't wealth or technology, it's the mastery of risk through probability – the shift from a world where outcomes were attributed to fate and divine punishment to a world where we could actually calculate and shape odds. It's humbling, honestly, to realize how recent that shift is in human history, and it's a good reminder every time I catch myself getting overconfident about a forecast.

For the purely mechanical, rules-based end of investing, there's Joel Greenblatt's *The Little Book That Beats the Market* – his Magic Formula, which is really just "buy good companies at bargain prices, consistently and mechanically," built around high return on capital and a cheap earnings yield. I don't run a mechanical screen personally, but I think the underlying discipline is correct: most investors underperform not because they lack intelligence, but because they chase hot stocks, overreact to news, and abandon good strategies the moment they underperform for a quarter or two. Discipline beats genius more often than genius beats discipline.

And then there's Saurabh Mukherjea's *The Unusual Billionaires*, which is close to home because it studies Indian companies specifically – Asian Paints, HDFC Bank, Nestlé India, Marico, and – this is a fun coincidence given what's in our research this quarter – Page Industries. His central question is why a handful of companies in a volatile, governance-challenged emerging market manage to generate sustained twenty-percent-plus returns on capital for decades, while most don't, and his answer is that it's never about political favoritism or leverage or riding a commodity cycle. It's disciplined capital allocation and governance, compounded patiently, for a very long time. That's basically the entire aspiration for this eight-stock portfolio in one sentence.

Now, here's a genuinely funny bit of overlap I only noticed while putting this quarter's research together – Max Gunther's *The Zurich Axioms* and *How to Get Lucky*. The Zurich Axioms are literally the risk philosophy of old Swiss bankers and traders – the idea that uncertainty isn't a problem to be engineered away, it's the natural condition of being alive, and that wealth comes from learning to live intelligently with risk rather than from avoiding it entirely. Gunther's line that if you're not worried, you're probably not risking enough has stuck with me. And funnily enough, this quarter's macro research happened to include a full deep-dive on Switzerland itself – so there was a strange, pleasant symmetry in reading about the psychology of Swiss speculation in the same month I was reading the actual balance sheet of the Swiss Confederation. More on that in a bit.

The rest of the shelf is less directly about markets, but I don't think it's separate from this work either. Peter Attia's *Outlive* is about longevity – and I think a lot about my own health as literally the compounding asset with the longest duration I own, longer than any stock. If I want to actually watch a multi-decade thesis play out, I need to still be around and thinking clearly in decade three and four. Christopher Hitchens' *God Is Not Great* isn't about investing at all, but the habit of mind it teaches – question the story you're told, especially the comforting one, especially the one everyone around you already believes – is a discipline I try to apply everywhere, markets included. Bill Bryson's *A Short History of Nearly Everything* keeps me humble about how much of the physical world I genuinely don't understand, which I think is healthy medicine against the kind of overconfidence Taleb spends three books warning about. And yes, John Gray's *Men Are from Mars, Women Are from Venus* is on the shelf too – not everything I read is about capital allocation, and I don't think it should be. A mind that only reads finance books tends to produce pretty narrow investing.

If there's a single thread running through all of it – Taleb's randomness, Marks's cycles, Dalio and Sharma's nations, Bevelin's mental models, Damodaran's narrative-and-numbers, Mukherjee's patient compounders – it's this: be humble about what you can predict, be honest about what's luck versus skill, think in decades rather than quarters, and never mistake a good story for a good number, or a good number for a good story. You need both. That's the lens everything below gets filtered through.

On AI, India, and Why I'm Not Panicking About Either

Let me talk about the macro picture for a minute, because I get asked about this constantly, and I want to give you my actual view rather than a hedge.

On AI – I think it's genuinely groundbreaking, and I think the world a few years from now looks meaningfully different from the world today because of it. But I want to separate two questions that people constantly collapse into one: is AI transformative, and is AI something India specifically needs to be racing to build at the frontier? My answer to the first is yes. My answer to the second is no, not right now, and I think conflating the two leads to bad capital allocation decisions.

Here's the way I think about GDP growth at the country level – and this connects directly to the Dalio and Sharma frameworks I just walked you through. Growth comes from a handful of variants: productivity growth, labor force growth, labor education and skill quality, and capital deepening – the accumulation of machinery, infrastructure, and tools per worker. Every country has to choose where it pushes hardest, because you can't maximize all of them simultaneously with limited state capacity and limited capital.

My view is that India doesn't need to push on frontier technology and productivity right now. That race is being run by the US, South Korea, and Taiwan, and there will almost certainly be a spillover effect – a purge-and-diffuse pattern, where the frontier innovation eventually filters down and becomes cheaper and more accessible to economies like ours, the way it always has historically. Chasing the frontier directly, before the low-hanging fruit is picked, is expensive and often wasteful.

What India actually needs, in my view, is to focus on the boring stuff – and I mean that as a compliment, not a criticism. Railways. Education infrastructure. Transportation. Ease of doing business. More competitive, more predictable regulatory environments. Give the population and the businesses genuine societal support to experiment, to start things, to fail cheaply and try again. The productivity-frontier gain – the AI-driven leap

– is still a fair distance away for an economy at India's stage. But there's a full decade of GDP growth available just from clearing out the low-hanging fruit, the unglamorous structural stuff that doesn't make headlines but compounds relentlessly if you stay disciplined about it. That's the Sharma lesson again, really – reform momentum, sustained, on the boring things, beats a flashy bet on the frontier that the country isn't structurally ready to win.

Now, the question I get asked most often on this exact topic is job losses – specifically, whether AI is going to gut the Indian IT services sector, given that TCS sits at over 12% of this portfolio. I want to be straightforward: I don't believe that's the base case, and I remain constructive on Indian IT from here. My reasoning is fairly simple – this is, at its core, another mathematical modeling problem, another wave of automation layered on top of decades of prior automation waves that the sector has already absorbed and adapted around. Things adjust. The nature of the work shifts, the skill mix shifts, but I don't see a clean, sudden discontinuity where the entire sector's value proposition evaporates. I hold that view with appropriate humility – Taleb would tell me to be suspicious of any confident forecast about a genuinely uncertain, fast-moving technology shift, and I try to keep that voice in my head even while I explain why I'm not panicking.

So that's the macro backdrop I'm operating against – bullish on India's ability to grind out a full decade of solid, unglamorous growth; skeptical that the country needs to sprint for the AI frontier to get there; and comfortable, not complacent, about the IT sector's ability to adapt rather than collapse.

Choosing the Soil Before Choosing the Seed

Before I get into the actual research, let me remind you – and honestly, remind myself – of the framework this whole operation is built on, because it explains why two entire countries show up in a portfolio letter about eight Indian stocks.

The mechanism I use is what I call the Sovereign Bet. The idea, borrowed straight out of Darwin by way of Bevelin's book, is that you don't evaluate an organism in isolation from its ecosystem. A company doesn't compound in a vacuum – it compounds inside a country, inside that country's institutions, capital markets, rule of law, currency stability, and demographic trajectory. So step one, always, is choosing the soil. Only after I'm convinced the soil is worth planting in do I go looking for the specific organism – the company – worth planting there.

That's why, even though every single position in this portfolio is Indian, our research time this quarter went into studying two countries that aren't India at all: Switzerland and Australia. Not because I'm about to go buy Swiss or Australian equities next quarter – I'm not signaling that – but because understanding what genuinely stable, well-run soil looks like, in two very different flavors, sharpens the standard I hold India to. You can't know if your own backyard's soil is good if you've never actually walked through anyone else's.

The Soil, Case One: Switzerland

Let's start with Switzerland, because it's the cleanest possible case study in what Dalio would call a country deep in the disciplined, unglamorous stage of its internal cycle – and what Gunther, in the Zurich Axioms, would recognize as the psychological home turf of calculated risk-taking.

The headline growth number for 2026 isn't exciting on its own – real GDP growth is projected at somewhere between 1.0% and 1.4%, below-average by global standards, weighed down by external headwinds like US tariff shocks and a soft Eurozone. If you were doing first-level thinking, you'd stop right there and call it boring. But that's exactly the trap Marks warns about – the boring number isn't the story.

The actual story is structural. Switzerland ranks second globally in economic complexity, meaning its production base specializes in things that are genuinely difficult for competitors to replicate – precision instruments, advanced pharmaceuticals, high-end manufacturing. It's ranked first in the Global Innovation Index, and has held that position consistently since 2015. The country runs what's essentially a constitutional "debt brake" – a fiscal rule that keeps government debt as a share of GDP on a steadily declining path, sitting in the mid-thirties percent and still falling. Inflation has been consistently low, unemployment is low and falling, and the Swiss National Bank has genuine monetary sovereignty – the ability to decouple its own inflation trajectory from the Eurozone and the US, which is a rare and valuable kind of independence in a fragmenting world.

None of this is an accident, and none of it is recent. It's built on what I'd call institutional continuity – direct democracy and a consociational governing style that makes policy shift incrementally rather than violently, which means the risk of a sudden regime change or a populist rupture is structurally low. That's the Harari-and-Diamond lens again – deep institutional inheritance, not a lucky recent streak.

The honest risks are real too, and I want to give them equal airtime rather than just cheerleading. Switzerland runs one of the highest cost bases in the world for labor and operations, which means it has to keep out-innovating its own cost structure just to stay competitive – there's no slack to coast on. And it has real asymmetric exposure to external tariff regimes, particularly from the United States, given how export-dependent the economy is. This is a country that has to keep earning its stability every year; it's not a birthright that runs on autopilot.

My honest takeaway, filtered through Dalio: Switzerland reads like a great power in the mature, disciplined phase of its internal cycle – sound finances, high productivity, strong institutions – rather than the early, hungrier phase, or the late, complacent one. It's not a country I expect explosive growth from. It's a country I'd point to as the reference standard for what "safe, predictable, multi-decade compounding soil" actually looks like, so that when I evaluate India's own institutional trajectory, I have something real to measure it against.

The Soil, Case Two: Australia

Australia is a genuinely different flavor of stable soil, and I think the contrast with Switzerland is more useful than either country studied alone.

Where Switzerland compounds on precision and complexity, Australia compounds on something almost opposite – it's a commodity-intensive, export-led economy sitting on top of what's colloquially described as a periodic table that lights up like a Christmas tree: iron ore, coal, liquefied natural gas, critical minerals, all funding a high-standard-of-living, services-driven domestic economy. That's a structurally different bet than Switzerland's – Australia's prosperity is tied much more directly to global commodity demand and to its relationship with its biggest trading partners, which makes it a genuinely different risk profile even though both countries carry top-tier credit ratings.

Institutionally, Australia is strong – an independent central bank, a resilient judicial system, and a top-tier position on global rule-of-law rankings. Sovereign debt is low and manageable relative to growth, and the AAA credit rating looks secure on the forward numbers. There's also a genuinely interesting demographic engine here: a "Skill Stream" migration policy that keeps bringing in human capital on a continuous basis, which is a real structural advantage most developed economies don't have.

But — and this is the honest Sharma-style caveat — productivity, measured as total factor productivity, is close to flat. Near zero. That's the single biggest risk to the multi-decade compounding thesis for Australia, because you can bring in as much labor and capital as you want, but if you're not getting more output per unit of input over time, growth becomes a story about inputs rather than genuine value creation. Layer onto that a housing market where the average home price now sits above a million Australian dollars, household debt levels that are high by global standards, and a current account that's been widening rather than narrowing, and you get a picture of an economy with real structural comfort but a productivity problem it hasn't solved.

The more encouraging piece is the recent policy pivot — a "Future Made in Australia" industrial strategy that's starting to direct capital toward genuinely high-complexity sectors, critical minerals processing and green energy infrastructure in particular, rather than just extracting and shipping raw commodities. If that pivot sticks, it's the kind of reform-momentum story Sharma would recognize as a country trying to climb the complexity ladder before the current growth model runs out of road.

My honest read: Australia is exceptionally reliable soil in the institutional and sovereign-credit sense — genuinely one of the safest jurisdictions for long-term capital anywhere in the world — but it's carrying a productivity question mark that Switzerland simply doesn't have. It's a "friend-shoring" partner of choice in a fragmenting geopolitical world, sitting on a deep domestic savings pool, and I think it's well placed to navigate the next decade. But if I were underwriting it the way I underwrite an actual position rather than a reference case, the TFP stagnation would be the thing I'd want resolved before I got excited.

On My Desk This Quarter — The Watchlist

Now for the part that's genuinely new this quarter — the research itself, folded straight into this conversation instead of living in a separate document you'd have to go read on your own.

I want to be upfront about what this section is and isn't. This is not a list of new buys. Of the names I'm about to walk you through, exactly one — Caplin Point Laboratories — is actually sitting in the eight-stock portfolio today. Everything else here is what I'd call active reading — companies and, in a few cases, entire sectors I've spent real hours on this quarter, building out the same kind of business-understanding-to-conclusion framework I'd apply to something I already own, precisely so that if and when a position gets sold, I already know what I've earned the

right to replace it with. Munger's inversion principle applies here in a very literal way – instead of asking "what should I buy," I spend most of my research time asking "what would disqualify this," and only the names that survive that filter stay on the desk long enough to be worth telling you about.

So think of this less as a shopping list and more as a diary of what's been occupying my attention – the pattern-recognition training I'm putting myself through, quarter after quarter, so that when a genuine opportunity to add a ninth name does show up, I recognize it immediately instead of scrambling to understand it from scratch.

Let me start with the one that actually matters most to your capital right now, because it's already in the book.

Caplin Point Laboratories – The One We Actually Own

Caplin Point sits at 15% of the portfolio, and it's earned that weight the hard way – through what I'd call a genuinely high-conviction case of a business that de-risked itself deliberately rather than by accident.

The core of the thesis is geographic diversification combined with a balance sheet that's essentially free of leverage. Caplin's therapeutic footprint spans dozens of segments, and its historical playbook has been to target markets that bigger, more prestigious pharma companies consider too small or too unglamorous to bother with – underserved and unserved geographies, mostly across Latin America, where the company built a distribution model that gets essential, life-saving medication to roughly thirty thousand touchpoints without the overhead of a conventional multinational pharma sales machine. That's a very Mukherjee-style story, honestly – not a flashy growth narrative, just disciplined execution in a corner of the market nobody else wanted to fight for.

The company also benefits from a genuine cost advantage in how it sources its active pharmaceutical ingredients, and it expenses its R&D rather than capitalizing it – a small accounting detail that actually tells you something real about management's conservatism, because it means the reported earnings aren't being flattered by pushing costs onto the balance sheet. Revenue has grown from roughly ₹1,553 crore to ₹2,034 crore over the period we tracked, with gross margins in the high fifties to low sixty percent range – genuinely rare margin territory for a company operating in "bottom of the pyramid" markets rather than premium ones.

The company is currently in what it calls its "Caplin 2.0" phase, and there are three specific growth levers I'm watching closely. First, a push toward a "Made in USA" positioning for its American business, which eliminates profit-sharing arrangements and should be a genuine margin expander as more products launch under that label. Second, a move into oncology and other specialized, high-barrier verticals through ANDA acquisitions and a dedicated manufacturing site – that segment is expected to become a real profit driver starting around FY26/27. And third, a backward-integration push – internally called Project 90 – to build captive API supply for the more than ninety core products it sells in the US and globally, which matters a lot for injectable products where API cost is a large chunk of gross margin. Management is targeting \$100 million in US revenue by FY27/28 and aims to double LATAM revenue in the same window.

The honest risk here isn't leverage – there basically isn't any – it's execution risk on three growth initiatives running in parallel, and the ordinary regulatory risk that comes with any pharma manufacturer selling into the US market. But following a clean USFDA inspection in August 2024, the regulatory runway looks clear for now. This is the kind of business Taleb would call genuinely antifragile in its financial structure, even if the operating business still has to execute – zero leverage means a bad quarter doesn't threaten the company's survival, it just delays the thesis.

Dixon Technologies – The Brand Behind the Brands

Dixon is fascinating precisely because most people who use its products every day have never heard of it. It's an outsourced manufacturing partner – an Original Design Manufacturer, to use the technical term – for a huge array of global and domestic electronics companies, making everything from LED televisions and AC circuit boards to washing machines and mobile devices, so that the brand owners can focus purely on marketing and distribution while Dixon handles the capital-intensive, operationally messy business of actually building the things.

What I like about it is the direction of travel – it's moving from being a pure assembly house toward a more vertically integrated manufacturer, which is a structurally better place to sit in the value chain, and it's got real optionality in automotive electronics through a joint venture, betting on the idea that vehicles are becoming computers on wheels, which plays directly to Dixon's core competency in electronic sub-assemblies. Balance sheet is close to spotless, with debt-to-equity sitting in the 0.07 to 0.09 range.

What I don't like, and what keeps this off the actual portfolio for now, is the valuation and the concentration risk – a big chunk of mobile revenue, roughly 72%, comes from a single customer, and the stock trades at a premium that assumes a lot of that growth continues flawlessly. It's a name I keep reading, not one I'm ready to own yet.

There's also a Bevelin-style Darwinian question I keep sitting with on this one: is Dixon's moat the kind that adapts as the electronics value chain evolves, or is it a moat that's specifically well-suited to today's assembly-heavy landscape and could get disrupted the moment brand owners decide to bring manufacturing closer in-house, or automate it more aggressively than a contract manufacturer can. I don't have a confident answer yet, and until I do, this stays a name I read every quarter rather than one I act on.

The customer list is genuinely worth knowing, because it tells you how central Dixon has become to the physical reality of "Make in India" – it manufactures for Samsung, Xiaomi, Motorola, Panasonic, and Philips on the electronics side, and increasingly for HP, Asus, Acer, and Lenovo in IT hardware. That's not a company riding one relationship; it's infrastructure that half the electronics brands you actually use every day quietly depend on. Which is exactly why that adaptability question matters so much here – the moat is real today, but it's built on being the most convenient plug-and-play manufacturing partner in the country, and convenience-based moats have to keep re-earning themselves every contract cycle in a way an entrenched brand moat doesn't.

Apollo Hospitals – The Total Health Partner

Apollo is best understood not as a chain of hospital buildings but as an integrated healthcare ecosystem spanning the entire patient journey – 73 hospitals with roughly 10,200 beds for high-acuity care, a digital and pharmacy distribution arm called Apollo HealthCo running over 7,000 pharmacy stores, and a primary-care and diagnostics arm underneath that. It's the bridge between elite clinical care and retail accessibility, which is a genuinely hard thing to build.

The founder, Dr. Reddy, is still the visionary head in his nineties, with day-to-day operations run by his four daughters across clinical, financial, digital, and technology functions respectively – a genuinely unusual and, so far, well-functioning succession structure. The company's roadmap through 2030 involves adding several thousand more beds in high-yield metro clusters and potentially listing the HealthCo digital arm separately to unlock its value. For an investor, the pitch is fifteen-to-eighteen

percent revenue visibility with improving capital efficiency in a sector where the barriers to entry get higher every year. It's on the desk because Indian healthcare demand is a genuinely durable, multi-decade tailwind – I just haven't yet convinced myself the valuation and the complexity of the multi-entity structure make this the right entry point.

The Damodaran lens is particularly useful here, actually, because Apollo is a genuine case of a good story running slightly ahead of a fully resolved set of numbers – a potential HealthCo listing that would unlock value on paper, a succession structure that looks well-functioning today but hasn't yet been tested by a real generational handover, and bed-count expansion plans that sound compelling but haven't all converted into signed, funded projects yet. None of that makes it a bad business. It makes it a business where I want to see the narrative and the numbers converge a bit further before I'm willing to size a position around the story.

What actually convinced me the demand side is close to bulletproof is a single, slightly grim statistic – non-communicable diseases already account for 62% of deaths in India and are projected to climb to 84% by 2030, and Apollo's most complex specialties, what the company internally groups as its "CONGO" divisions – cardiac sciences, oncology, neurosciences, gastroenterology, and orthopaedics – sit directly in the path of that trend. Nobody delays a cardiac bypass the way they'd delay a car purchase in a weak quarter, which is exactly the kind of non-discretionary demand Apollo Hospitals as a category benefits from regardless of the broader economic cycle. And the digital transformation has genuinely happened, not just been announced – Apollo 24/7 has gone from effectively nothing five years ago to over 36 million registered users and roughly 14,500 daily virtual consultations today, which is the kind of platform-style scaling that, if HealthCo ever does get separately listed, could justify a very different multiple than a pure hospital-bed operator would command.

DMart – The Volume Machine

DMart, run under Avenue Supermarts, is Radhakishan Damani's disciplined answer to Indian retail – a value-retailing supermarket chain that wins not through marketing or prime real estate but through relentless cost discipline passed straight through to the customer. It's the kind of business Greenblatt would appreciate immediately – high returns on capital, earned the boring way.

With only 442 stores in a country of 1.4 billion people, the expansion runway is enormous, and the company reinvests

essentially all its profits into new stores rather than ever touching the debt markets – a genuinely rare discipline in Indian retail. The real tension worth watching is digital: DMart Ready, its delivery arm, has to compete with quick-commerce players on convenience without abandoning the everyday-low-price positioning that built the brand in physical stores. Private labels, currently fifteen to twenty-three percent of sales, are the next real margin lever if that share keeps climbing toward thirty percent. This is exactly the kind of "boring compounder" Mukherjea's book is built around, and it stays firmly on my radar.

What keeps me from acting on it yet is almost entirely a valuation question rather than a business-quality one – the market has never really let DMart trade cheap, precisely because everyone else has read the same Mukherjea-style compounder thesis I just gave you. That's the Marks second-level-thinking trap in its purest form: the consensus view on DMart being an excellent business is correct, which is exactly why it's hard to find genuine mispricing in it. I'd rather wait for a moment when the market's patience with the stock's rich multiple runs thinner than its patience with the underlying business, than force my way into a name at a price that already assumes everything goes right.

What I find genuinely elegant about the model is how un-discretionary it is by design – more than 77% of what DMart sells is the kind of everyday household staple that gets bought in good years and bad years alike, groceries and dairy and home care, not the apparel and appliances that make up the smaller, more cyclical slice of the mix. Damani built "Everyday Low Price" as a consistent, year-round promise rather than a calendar of sale events, and that consistency is itself the moat – customers stop comparison-shopping once they trust the baseline price is already the lowest they'll find. It's a very Adam Smith kind of insight, honestly: specialize ruthlessly in doing one unglamorous thing – low-cost distribution of daily necessities – better than anyone else, and let that specialization compound.

Maruti Suzuki – The Bellwether at a Turning Point

Maruti is forty-three years into being the dominant force in Indian passenger vehicles, majority-owned by Japan's Suzuki Motor Corporation, and it's now repositioning itself for what the company calls its "Vision 3.0" era – an aggressive push into SUVs, where it's already reclaiming meaningful market share, alongside a genuinely pragmatic multi-pathway strategy across CNG, hybrids, and EVs rather than betting the company on a single powertrain technology.

Financially it's close to a textbook example of what Taleb would call robust rather than fragile – zero debt, return on capital employed above twenty percent, and enough working-capital discipline to fund a seventy-thousand-crore expansion plan without diluting shareholders. Competition from Mahindra and Tata in the SUV and EV segments is real and intensifying, but I think the multi-pathway approach is the correct hedge against the risk of picking the wrong technology winner too early. Watching this one closely as the SUV transition plays out over the next couple of years.

I'll admit this is one where the Taleb-versus-Damodaran tension in my own head is most visible. The Damodaran instinct says pick a scenario, run the numbers, commit to a view on whether the SUV pivot succeeds. The Taleb instinct says the actual value of Maruti's strategy is precisely that it doesn't require picking a winner among CNG, hybrid, and EV – it's a barbell across power-train outcomes rather than a single confident bet, and the zero-debt balance sheet means the company can afford to be wrong about the pace of any one technology without it becoming an existential problem. I've come down on the side of thinking the optionality itself is the asset here, which is part of why this stays high on the watchlist even without a fully resolved single-scenario valuation model.

Suzuki Motor Corporation holds a 56.28% stake, which matters because it means the parent's capital and technology roadmap – including the EV platforms being developed globally – flow fairly directly into the Indian entity rather than staying locked up overseas. The product ladder is genuinely comprehensive, from the Alto K10 at the entry-level, first-time-buyer end all the way up through the Baleno and Swift for middle-class urban families to the Grand Vitara, Invicto, and Jimny higher up the SUV range, and the company has quietly built real annuity-like businesses around the core vehicle sale – Genuine Parts and Accessories alone contributes close to 12.78% of turnover, and the "True Value" pre-owned network and "Smart Finance" arm mean Maruti captures value across the entire ownership lifecycle, not just at the point of first sale. That full-lifecycle capture is underrated, I think, relative to how much attention the monthly sales-volume headlines get.

State Bank of India – The Path to Ten Trillion

SBI is essentially the financial nervous system of the country – a universal banking conglomerate mobilizing savings from over fifty crore customers and deploying that capital everywhere from rural micro-enterprises to national infrastructure, while cross-selling insurance, mutual funds, and investment banking through its

subsidiary ecosystem. It's evolved from what people still picture as a sleepy legacy public-sector lender into what the research calls a genuine "financial supermarket."

The numbers are the interesting part – return on assets around 1.10% and return on equity approaching twenty percent, which is private-sector-grade profitability from a state-owned institution, alongside a track record of cleaning up asset quality that was genuinely troubled a decade ago. As India moves toward a ten-trillion-dollar economy, SBI's moat – built on physical reach, trust, and an increasingly capable digital layer – looks more durable than the market's historical skepticism of PSU banks would suggest. It's the kind of "boring but essential" financial infrastructure name I keep circling back to.

The "Lindy Effect" argument for SBI is honestly hard to ignore once you see the numbers – this is an institution dating back to 1806, with over 22,500 branches and 63,000 ATMs giving it a physical convenience moat in semi-urban and rural India that no digital-first private competitor can replicate quickly, and it manages roughly 15 crore accounts opened under the government's Jan Dhan financial-inclusion push alone. What's genuinely changed is the digital layer on top of that legacy footprint – the YONO platform has turned a paper-heavy legacy lender into something that can compete with private-sector apps on convenience while still owning the trust and reach that two centuries of sovereign backing built. That combination of old-world reach and new-world interface is a harder moat to disrupt than most people give it credit for.

IREDA and REC – The Twin Financiers of India's Power Buildout

I looked at these two together because they solve a similar problem from slightly different angles. IREDA is a "Navratna" government enterprise that's essentially a pure-play financier of the renewable energy transition – solar, wind, green hydrogen, and increasingly, retail lending for rooftop solar and EV infrastructure. What stands out is the credit quality trajectory – non-performing assets falling from 8.8% down to 3.75% over eighteen quarters, which suggests the underwriting discipline has genuinely improved rather than just gotten lucky with a benign cycle. The main risk is concentration – a single large project defaulting can move the needle on asset quality given how wholesale-heavy the loan book is.

REC, formerly Rural Electrification Corporation, is the older, larger sibling – a Maharatna enterprise financing power generation, transmission, and distribution at scale, now also diversifying into logistics and broader infrastructure lending.

It's cleaned up meaningfully since the 2015-to-2018 NPA crisis in the power sector, with net NPAs now down around 0.3 to 0.4%, and it's compounding toward a ten-lakh-crore loan book at a consistently high return on equity near 21%. Both of these read like quasi-sovereign proxies on India's energy build-out, and I think of them as a pair rather than picking one over the other.

Both institutions exist to solve the same underlying market failure – commercial banks fund themselves with short-term deposits, which makes them structurally uncomfortable lending for twenty-five-year power infrastructure, so IREDA and REC step in as patient, sector-specialist capital instead. IREDA, founded to promote renewables specifically, positions itself as a "knowledge partner" as much as a lender, offering technical consultancy alongside project financing. REC, dating back to 1969 and carrying Infrastructure Finance Company status, leans even more heavily on its sovereign-adjacent credit rating, with roughly 87 to 89% of its loan book still concentrated in State Electricity Boards and public power utilities, letting it borrow cheap in global capital markets and re-lend into infrastructure at tenures – up to 25 years – that no ordinary commercial bank could comfortably match. That's the whole game for both of them: borrow long-duration trust, lend it into projects nobody else has the balance sheet patience for.

The Greenblatt-style discipline is genuinely useful on both of these – strip away the "green energy" narrative excitement for a moment and just look at return on equity against asset quality, the same way you'd screen any lender. On that basis, both clear a fairly high bar, which tells me the appeal here isn't purely thematic enthusiasm about the energy transition, it's underwritten by actual credit discipline. That distinction matters to me, because I've watched enough "good story, mediocre balance sheet" businesses get re-rated purely on narrative to be suspicious of my own excitement whenever a thematic tailwind and a decent business happen to coincide.

DLF – From Leveraged Developer to Cash Machine

DLF is the story of Indian real estate maturing in real time. It's the dominant institutional player in the space, most visible through Cyber City in Gurugram, and it runs two engines – a development business selling residential and super-luxury units, and an annuity business, run through a joint venture with Singapore's GIC, leasing out roughly forty-nine million square feet of office and retail space.

What makes this one interesting now rather than five years ago is the balance sheet transformation – DLF has gone from a heavily

leveraged developer to a business sitting on an eleven-thousand-six-hundred-crore cash pile, funding its own growth without relying on external credit. The risks are concentration in the Gurugram market and sensitivity to the luxury demand cycle, but the pivot toward managed townships, senior living, and data centers signals forward-looking management rather than a company just riding the current residential upcycle. A genuine proxy for the formalization of Indian real estate, and worth tracking even without an obvious near-term entry point.

The internal split between what the company calls its Devco and Rentco engines is a genuinely useful mental model even outside real estate – Devco sells, Rentco leases, and the leased side, run through the DCCDL joint venture with Singapore's GIC, means DLF isn't just exposed to Indian residential sentiment, it's also a landlord to Fortune 500 multinationals and Global Capability Centers renting nearly 49 million square feet of Grade A office space. That tenant base is itself a quiet bet on India's role as a back-office and innovation hub for global companies, which is a different and arguably steadier demand driver than the luxury-apartment cycle that gets most of the attention.

Page Industries – The Jockey and Speedo Licensee

Page Industries is actually one of the companies Mukherjee studies directly in *The Unusual Billionaires*, which made revisiting it this quarter feel like closing a loop. As the exclusive licensee for Jockey and Speedo in India, it's built a business on a genuinely simple insight – comfort is a basic, recurring human need, and Page turned innerwear and athleisure from a functional afterthought into a lifestyle category.

The numbers back up the reputation – return on capital employed above 70%, with essentially no financial leverage. The near-term challenge is premium competitors like Van Heusen encroaching from above and the general shift toward digital shopping, but the structural moats – brand stickiness, manufacturing depth, an enormous distribution network – look intact, and a new plant in Odisha plus expansion into the Gulf region give it real runway. A new CEO, promoted from within and focused on digital and growth, is the thing I'd want to watch play out before forming a stronger view either way.

Worth knowing just how concentrated this business actually is underneath the brand-stewardship story – Jockey alone contributes roughly 99% of revenue, with Speedo, despite leading the niche swimwear category, making up the remaining 1%. Since the company's founding in 1994, the real achievement has been almost sociological rather than purely commercial: taking innerwear "out

of the closet," in the company's own phrase, through international-style retail fixtures and box packaging, and turning a functional necessity into a category people would proudly display in an organized store. Page currently commands an 18-to-20% penetration rate in its primary high-income and upper-mid-income segments, which tells you the brand-building worked – the open question, and it's the one I'm watching the new CEO on, is whether that same playbook translates as cleanly to a digital-first shopping experience as it did to the physical retail environment it was originally designed for.

The Power-Grid Cluster – Hitachi Energy, BHEL, and GE Vernova T&D

I spent a genuinely large chunk of this quarter's research on the electrification theme, because I think it's one of the more durable, decade-plus structural stories in industrials right now, and three names anchor it.

Hitachi Energy is essentially the nervous system of the modern power grid – the infrastructure that lets electricity move from wind farms, solar arrays, and traditional plants to cities, industrial complexes, and the data centers powering the AI buildout. The business is asset-heavy but consistently outperforms peers on returns, and it's got a genuinely interesting long-term optionality – as the roughly half a million grid assets it has in the field generate more operational data over time, there's a real path toward a "grid-as-a-service" model that could re-rate the business from an industrial multiple toward something closer to a tech-infrastructure multiple. The risk is pure execution – scaling manufacturing and talent fast enough to meet 2030-level demand – but with nine billion dollars of committed investment behind it, the ambition is at least funded.

The actual hardware is worth naming because it's more specialized than "grid equipment" suggests – high-voltage direct current systems capable of moving huge amounts of electricity over long distances with minimal loss, the full range of transformers from small distribution units up to ultra-high-voltage systems, and gas- and air-insulated switchgear. Layered on top is genuinely interesting software, including an AI diagnostic platform called HMAX Energy that predicts equipment failure before it happens. That combination of specialized, hard-to-replicate hardware and a growing software layer is exactly the kind of moat that tends to compound rather than commoditize over time.

BHEL is the older, more state-anchored piece of this story – India's foundational heavy-engineering player, building the turbines, boilers, and generators that keep the national grid running, while also serving as a manufacturing arm for defense and

rail. It's in the middle of a genuine resurgence, executing against a two-trillion-rupee order backlog, pivoting away from its historical coal dependence toward a more diversified base spanning renewables, rail, and defense. The thesis here is straightforwardly operational leverage – as that backlog converts to revenue, profitability should expand meaningfully, assuming execution holds.

GE Vernova T&D India is the third leg – a more than century-old presence in Indian power transmission and distribution, now the listed vehicle for GE Vernova's grid solutions business globally. This one has genuinely excellent numbers: debt-free, return on capital employed near 55%, and an order backlog worth two and a half times trailing revenue. The catch is the valuation already reflects a lot of that excellence – trading in the eighty-to-ninety times earnings range – so the market has clearly noticed this story already. The risk isn't the business quality, it's paying too much for quality that's already been discovered, which is exactly the kind of trap Marks would flag: everyone agrees it's a good business, so the question becomes whether that consensus is already fully priced.

Taken as a trio, Hitachi Energy, BHEL, and GE Vernova T&D are also a nice illustration of Sharma's reform-cycle logic operating at the sector level rather than the country level – a genuinely underinvested piece of infrastructure, the grid, suddenly getting a decade's worth of catch-up capital thrown at it all at once because electrification and AI-driven power demand made the old grid's limitations impossible to ignore any longer. The risk with any theme like this is timing – being early to a real structural story is functionally indistinguishable from being wrong, for however long the "early" part lasts, and grid infrastructure build-outs have a well-documented history of taking longer than initial announcements suggest. I'd rather stay a close reader of this cluster for another few quarters than rush into whichever name currently has the cleanest narrative.

The Pharma-and-Water Cluster – CDMO, Denta Water, and the Broader Water Sector

Two sector-level deep-dives rounded out the quarter, both structurally interesting even without a specific name I'm ready to act on yet.

On the pharmaceutical manufacturing side, I spent time on the broader Contract Development and Manufacturing Organization theme – the shift in Indian pharma from being purely the "Pharmacy of the World," selling affordable generics, toward becoming what the research calls the "Laboratory of the World," doing higher-value

contract manufacturing and even biosimilar development for global players. The industry's moving from oral solids toward the much harder physics of sterile injectables and biosimilars, and there's a real "Lindy Effect" at work – regulators trust manufacturing facilities more the longer they've maintained a clean compliance record, which becomes its own kind of moat over time. Geopolitical tailwinds, particularly friend-shoring away from Chinese API dependence, could let Indian CDMOs meaningfully grow market share by 2028.

The scale underneath that story is genuinely large already – India supplies roughly 20% of the world's generic medicine exports and about 40% of the entire United States' generic drug demand, on top of a domestic pharmaceutical market worth around 60 billion dollars and global exports near 30.5 billion dollars in the last fiscal year, with the sector projected to keep compounding at a 10-to-12% annual clip through 2030, comfortably outpacing GDP growth. This is Caplin Point's actual neighborhood, in other words – the CDMO shift toward sterile injectables and biosimilars is the more capital-intensive, harder-to-enter version of the same broad industry Caplin has been quietly compounding inside for years.

On water, India is facing a structural supply-demand gap serious enough that unmanaged, it could shave a real percentage off GDP by 2030, and government programs like the Jal Jeevan Mission are pushing capital into both traditional pipe infrastructure and more sophisticated treatment technology – desalination, advanced filtration, zero-liquid-discharge systems. Denta Water and Infra Solutions is the specific name I looked at here – a Karnataka-based company that pivoted from general real estate into a genuinely specialized niche, groundwater recharge using recycled water, running an asset-light model with a virtually debt-free balance sheet and strong margins. The catch, and it's a real one, is concentration – most of its revenue sits in Karnataka with a handful of government clients, which makes it a fundamentally binary bet on continued state-level policy support. It's expanding into other water-stressed states and even into transport infrastructure as a diversification hedge, but I'd want to see that geographic spread actually materialize before this becomes anything more than a name I'm watching.

The Plumbing Underneath Everything – Banking, NBFCs, and Life Insurance

The last piece of this quarter's research wasn't about any single company at all – it was about rebuilding the sector-level frameworks I use to judge every financial holding in this portfolio, because if you look at the eight names again, four of them – Bajaj Finance, Cholamandalam, Muthoot Finance, and Aptus –

are lenders of one kind or another, and I wanted to make sure the underlying model I'm judging them against was current.

On banking, the framework comes back to something simple: a bank resolves the friction between depositors who want liquidity today and borrowers who want long-term, illiquid funding for years – maturity transformation, in the technical language. The switching costs that create a real moat are structural, not sentimental – deep integration into a corporate client's credit and cash-management systems, or a retail customer's entire financial life, is genuinely hard to unwind, which is why exit barriers in banking run so high once a relationship is established.

On NBFCs – the space where Bajaj Finance, Cholamandalam, and Muthoot actually live – the founding purpose is filling the gap banks structurally can't fill, because banks are bound by conservative risk appetites and can't easily underwrite borrowers without clean, formal income documentation: the self-employed, unorganized retail, small and micro enterprises. NBFCs built relationship-driven, feet-on-the-street credit models to serve exactly that population, and increasingly, banks treat them as efficient distribution partners rather than pure competitors, routing indirect lending into rural and semi-urban markets through them. NBFC returns on equity in India average somewhere in the fifteen-to-twenty-percent range, but the sector runs on wholesale funding rather than cheap retail deposits, which means the whole model is more sensitive to funding-cost cycles than a deposit-funded bank is – a genuinely important distinction when I'm sizing conviction in Bajaj Finance versus SBI, for instance.

And on life insurance – not currently a portfolio holding, but a sector I think is worth understanding given how it interacts with household savings flows more broadly – the industry is shifting away from being a conservative, tax-incentivized savings vehicle toward a more competitive, protection-first, digitally-integrated market. A younger, more financially literate demographic entering the workforce, combined with a tax regime that's steadily eroding the old Section 80C savings incentive, is forcing insurers to sell the actual value of protection rather than the tax deduction that used to do the selling for them. Good sector to keep half an eye on, even without a specific name that's earned a place on the desk yet.

I put all three of these together at the end deliberately, because I think this is the least glamorous research I did all quarter, and also probably the most useful – understanding the plumbing underneath four of my eight positions matters more than finding a shiny ninth one.

How I Actually Size a Position

Let me pull back from individual names for a moment and talk about mechanics, because I think this is the part people ask about least and probably need most – not which stock, but how much of it, and why that number and not a different one.

I don't use a single mechanical formula, and I want to be honest that anyone who tells you they do is probably oversimplifying their own process after the fact. But there is a real logic underneath the weights you saw earlier – Caplin at 15%, Chola at 14.33%, Bajaj at 12.68%, and so on down to KPIT at 8.41%. Three things drive that sizing.

The first is conviction, in the Damodaran sense – how well does the numbers-side of the story hold up against the narrative-side, and how much of a gap is there between what I think the business is worth and what I'm paying for it today. A wider gap, held with genuine conviction, earns a bigger position. A thinner gap, even in a business I love, gets a smaller one, because the margin of safety just isn't as generous.

The second is what I'd call correlation risk, borrowed loosely from Taleb's obsession with tail events. Four of my eight names are lenders. If Indian credit conditions genuinely deteriorated – a real recession, a funding crunch, a regulatory shock to NBFC risk weights – those four positions wouldn't move independently of each other, they'd move together, hard. So even though each individual lender might deserve a large position on its own merits, I have to size the group with the awareness that I'm really running one concentrated bet on Indian credit quality wearing four different tickers. That's part of why TCS and KPIT sit in the book too – genuine diversification away from that single macro factor, not just "software is a good sector."

The third is what Gunther would call the worry test from the Zurich Axioms – if a position's size doesn't make me pay attention, it's probably too small to matter, and if its size keeps me up at night for reasons that aren't about the business itself, it's probably too large. Fifteen percent in a single small-cap pharma name is a real bet. It should feel like one. If it didn't feel like anything, I'd worry I wasn't being honest with myself about what concentration actually means.

None of this is a formula you could code up and run. It's closer to what Bevelin would call structured judgment – probabilistic, not mechanical, built by running the same set of questions against every name until the pattern of what deserves size versus what doesn't becomes something closer to instinct than arithmetic.

What Would Actually Make Me Exit a Position

People ask me far more often what would make me buy something than what would make me sell it, which I think is backwards, because the sell decision is where most investors actually lose money – either by holding on too long out of attachment, or by panicking out too early on noise. So let me tell you, honestly, what would move me to exit one of these eight.

It is almost never going to be a bad quarter. A single disappointing print, on its own, tells you very little – Taleb's whole point about randomness applies just as much to short-term underperformance as it does to short-term outperformance. If Caplin Point missed a quarter because of a shipping delay in Latin America, that's noise, not signal, and reacting to it would be a mistake.

What would actually move me is a change in the story that Damodaran's framework is built around – if the narrative that justified the position stops translating into the numbers it was supposed to translate into, not for one quarter, but as a pattern. If Dixon's customer concentration stopped being an acceptable risk because the actual customer relationship deteriorated. If a lender's credit quality started degrading in a way that wasn't explained by the broader cycle, but by something specific to their underwriting getting sloppy. That's the kind of change that would move me, because it's evidence the story itself was wrong, not evidence that the market is temporarily mispricing a story that's still intact.

The second trigger is what Munger and Bevelin would call a change in the ecosystem rather than the organism – remember the Sovereign Bet framework. If India's own soil quality changed in some structural way – a genuine deterioration in the regulatory environment, a shock to the rule-of-law framework, something at the Dalio-cycle level rather than the news-cycle level – that would matter more than anything happening at any single company, because it would mean the whole thesis underneath all eight positions needed re-examining, not just one of them.

And the third, honestly, is opportunity cost – the Greenblatt-style discipline of recognizing when a name on my watchlist has become a genuinely better risk-adjusted bet than something already in the book. That's a harder trigger to act on than it sounds, because it requires admitting that a position I've already committed capital and conviction to might not deserve its spot anymore, and that kind of admission fights against a very human instinct to stay loyal to your own past decisions. I try to remind myself that the market doesn't know or care what I paid, or

how long I've held something, or how much research went into the original thesis. Only the forward-looking picture matters.

What I won't do is sell because a position is up a lot and "feels" expensive on vibes alone, or because it's down and I've lost patience. Both of those are emotional decisions dressed up as analytical ones, and Marks's whole second-level-thinking framework exists precisely to catch you doing that to yourself.

What the Actual Week Looks Like

Since I'm trying to preserve the thought process here rather than just the conclusions, let me walk you through what the research work underneath everything I've just told you actually looks like week to week, because I think the process matters as much as the output, and it's the part a polished quarterly letter usually hides.

It starts with reading, and it starts wide before it narrows. I don't open a terminal and start screening for stocks. I open annual reports, concall transcripts, and increasingly, primary government and central-bank data – the kind of source Dalio and Sharma both lean on, because secondary commentary about a country's fiscal position is always a step removed from the actual numbers. For something like the Switzerland or Australia work this quarter, that meant going through SECO projections, IMF country data, and central bank statements directly rather than trusting a news summary of them.

Then comes what I'd honestly call the boring middle – building out the same repeatable structure for every business or country I look at, regardless of whether I already suspect I'll like it or not. Business understanding first: what does this thing actually do, in plain language, stripped of jargon. Then the numbers: growth, margins, returns on capital, leverage. Then governance and management – who's actually running this, how are they paid, do their incentives point the same direction as mine. Then risks, stated as honestly and specifically as I can manage, because a risk section that only lists generic disclaimers isn't actually doing its job. And only at the end, a conclusion – because I've found that if I let myself write the conclusion first, I start unconsciously cherry-picking evidence to support it, which is exactly the kind of narrative-fallacy trap Taleb spends a whole book warning about.

A lot of names don't survive that process, and that's fine – that's actually the point. Rejection is the default outcome of good research, not a failure of it. I probably read closely on twice as many companies as make it into anything I'd call a real

watchlist, and the watchlist itself is maybe five times larger than what ever earns an actual position. Munger's inversion logic again – I spend more energy trying to disqualify an idea than trying to fall in love with it, because falling in love with an idea is easy and disqualifying it properly is the actual skill.

The country-level work this quarter – Switzerland and Australia specifically – took longer than any individual company write-up, and I think that's appropriate given how foundational the "choose your soil first" step is supposed to be. It's slower, less immediately rewarding work than reading about a specific business with a specific stock price attached to it, but I've come to think of it as the highest-leverage research I do, because getting the soil wrong makes every downstream company decision wrong by default, no matter how good the individual business analysis is.

And then, every quarter, there's this – sitting down and actually narrating the last three months of that process out loud to you, which is its own useful discipline, because explaining a decision clearly to someone else is one of the fastest ways to discover whether you actually understand it yourself, or whether you've just been running on a plausible-sounding story you never fully interrogated.

The Lesson Behind Going From a Wider Book to Eight Names

I want to be honest about something rather than just presenting the move to concentration as a clean, confident upgrade, because I don't think it was purely that. I think for a while, I was running a wider portfolio than my actual conviction justified, and I think the honest reason wasn't sophistication – it was fear, dressed up in the language of prudence.

Early on, holding more names felt like the responsible thing to do. It felt like risk management. And there's a version of that instinct that's genuinely correct – Taleb would be the first to tell you that over-concentration in a domain you don't understand well is not bravery, it's just a faster way to go broke. But there's another version of "diversification" that isn't really about risk management at all – it's about not wanting to be wrong in a visible, attributable way. If you hold twenty-five names and three of them disappoint, you can point to the other twenty-two. If you hold eight and one disappoints, there's nowhere to hide, and you have to actually sit with having been wrong about something you said, out loud, with conviction. I think a meaningful part of my earlier, wider portfolio was unconsciously built to avoid that discomfort rather than to genuinely manage risk.

Munger's inversion principle is what eventually got me to see this clearly – instead of asking "what's the right number of stocks to hold," I started asking "what would it mean if I actually only had eight genuinely good ideas at any given time, and I was just refusing to admit that to myself." And the honest answer was uncomfortable: I probably only ever had six to ten names I truly understood at a level deep enough to hold through a bad quarter without flinching. Everything beyond that was padding – positions I owned because owning them felt safer than the alternative of admitting my circle of genuine understanding was smaller than my portfolio implied.

So the move to eight names wasn't really a move toward more risk. It was a move toward more honesty about the risk I was already carrying, just poorly disguised across more tickers. A business you've only shallowly researched, sized at three percent instead of zero percent, isn't actually safer than not owning it – it's just a smaller version of the same unexamined risk, and Taleb's silent-evidence point applies here too: you don't see the shallow research failing you quietly, position by small position, the way you'd see one concentrated bet failing you loudly.

I don't say any of this to be self-critical for its own sake – I think excessive self-flagellation is its own kind of unproductive noise. I say it because I think the most useful thing I can model for you, quarter after quarter, isn't a track record of never being wrong. It's an honest account of how the thinking actually evolved, including the parts where the earlier version of me was fooling himself a little. If preserving the actual thought process is the point of this recording, that has to include the parts where the thought process was flawed, not just the parts that make for a clean narrative in hindsight.

Questions I Keep Getting, Answered Properly for Once

Let me use the rest of this dialogue to actually answer the questions I get asked most often, properly, instead of giving the short version I usually give when someone catches me between meetings.

"Why only India? Why not diversify internationally?"

Partly this goes back to the Sovereign Bet framework – I'd rather deeply understand one soil than shallowly understand five. India is the ecosystem I can actually do primary research on, read the regulatory filings of, understand the cultural and consumer context of, without relying entirely on secondhand analysis. Sharma's whole argument about reform cycles and elite behavior is genuinely hard to assess from outside a country – you need

texture, not just data. That said, this quarter's Switzerland and Australia research isn't nothing – it's me actively building the muscle to eventually evaluate foreign soil with the same rigor, in case a genuinely compelling reason to deploy capital there ever shows up. I'm just not going to fake that rigor before it's actually there.

"Why so concentrated in financials? Isn't that just a sector bet?"

It's a fair challenge, and I don't think the honest answer is "no, it's not a sector bet" – it partially is. But I'd frame it more precisely as a bet on Indian credit penetration specifically, which I think is still early relative to where a ten-trillion-dollar economy will eventually need it to be. Four different lenders – Bajaj Finance, Cholamandalam, Muthoot, and Aptus – are each underwriting a genuinely different slice of that story: consumer and SME lending, vehicle and MSME finance, gold-backed lending, and affordable housing finance respectively. They're correlated in a systemic-shock scenario, which I addressed earlier, but they're not identical bets on the way there. I'd rather own four well-underwritten expressions of a theme I have real conviction in than force myself into an unrelated sector purely to make a pie chart look more balanced.

"You're 22. Why should anyone weight what you say the same as someone with thirty years in markets?"

Honestly, you probably shouldn't weight it the same, and I'd be suspicious of anyone my age who told you that you should. What I'd ask you to weight is the process, not the years – whether the reasoning holds up, whether the risk framework is coherent, whether I'm honest about what I don't know. Bernstein's whole history-of-risk book is a reminder that formal probability theory is only a few centuries old – expertise in this field has never purely been about tenure, it's been about who actually engages seriously with uncertainty rather than pretending it away. I'm not going to pretend three years of experience is thirty. I am going to argue that the quality of a decision doesn't automatically scale with the age of the person making it, and I'd rather be judged on the former.

"What happens if you're wrong about the India-doesn't-need-to-chase-AI-frontier thesis?"

This is the question I probably sit with the most, if I'm honest, because it's the one place in this letter where I'm making a genuinely contrarian, non-consensus call in the Marks sense, rather than just describing established facts about a company's balance sheet. If I'm wrong – if the AI productivity gap between

economies racing to the frontier and economies waiting for spillover effects turns out to be much larger and much more permanent than I'm currently modeling – then TCS and KPIT specifically would be more exposed than I currently believe, and I'd need to revisit that twenty-percent-plus software weighting. I don't think that's the likely outcome, but Taleb would tell me to always ask what I'd need to see to update, before I need to see it under pressure. What would update me: a sustained, multi-year divergence in productivity growth between frontier-AI economies and diffusion economies that shows up clearly in the hard data, not just in commentary. I haven't seen that yet. I'm watching for it.

"What's the actual difference between working at a family office and running an independent fund?"

This one comes up a lot, usually from people my age wondering whether they need to strike out entirely on their own to be taken seriously. I'd push back gently on that framing. A family office gives me something genuinely valuable that an independent setup at 22 almost certainly wouldn't – proximity to people who've already made the mistakes I'm currently trying to avoid making for the first time. There's a version of learning that only happens by sitting near someone who's lived through a genuine credit cycle, not just read about one in a Dalio book. I get that proximity daily, and I don't think I'd trade it for the independence of running something entirely my own a few years earlier than makes sense.

The honest tradeoff is that a family office comes with its own institutional preferences and constraints, and I'm not pretending this eight-stock portfolio operates in some kind of pure, unconstrained vacuum. But I think the framing of "family office versus independent" as a status ladder, where independent is inherently the more serious, more legitimate path, is itself a kind of narrative fallacy – a good story about ambition that doesn't actually track with where the best decision-making happens. Some of the most disciplined, longest-duration capital I've been exposed to sits inside family offices precisely because it isn't under quarterly redemption pressure the way a lot of independent funds are. That patience is worth more to me right now than the story of independence.

"Do you ever worry that reading this much philosophy and history is just procrastination dressed up as research?"

Honestly, yes, sometimes, and I think it's a fair thing to worry about. There's a real version of "reading widely" that's actually just an elaborate way of avoiding the harder, more uncomfortable

work of committing capital and living with the consequences. Bryson's book, in particular, is not going to help me size a position in Maruti Suzuki, and I'd be lying if I said every hour I spend reading is directly convertible into a better investment decision.

But I've made a kind of peace with it, and the peace looks like this: the books that read as pure procrastination in the moment tend to be exactly the ones that show up, unannounced, three months later, reframing a completely unrelated problem. The Harari and Diamond material on why certain institutions endure wasn't written with Switzerland's debt brake in mind, but it's the lens I actually used to think about why that debt brake has held for as long as it has. I can't always draw a straight line from a specific book to a specific decision in real time. I can tell you that, looking backward, the line is almost always there eventually. That's a harder thing to prove than a clean ROI calculation, and I hold it as a belief rather than a demonstrated fact, but it's a belief I'm willing to keep acting on.

"Why merge the research and the dialogue instead of keeping them separate, like before?"

Because I think the separation was, if I'm being honest with myself, partly a performance — it let the Research Notes sound authoritative and detached, and let the Dialogue sound warm and personal, and the truth is I'm the same person writing both, with the same uncertainties running underneath the confident-sounding research prose. I'd rather you hear the uncertainty where it actually lives, inside the research itself, than have it edited out to make the analysis sound cleaner than the thinking actually was.

The India Macro Backdrop, Concretely

Let me ground the "India will keep growing" view I mentioned earlier in the actual numbers on the table right now, rather than leaving it as a vibe, because I think a portfolio this weighted toward Indian financials deserves that specificity.

The RBI's Monetary Policy Committee has held the repo rate at 5.25% through its recent meetings, maintaining a neutral stance rather than signaling further cuts or hikes — which, after a run of cuts through late 2025, tells you the central bank thinks the current calibration is roughly right rather than still playing catch-up. Inflation has genuinely surprised to the downside — a long, nine-month stretch of falling food prices pulled headline CPI down to an eight-year low around the middle of last year, comfortably inside the RBI's target band, helped along by the GST

rate rationalisation that kicked in through September 2025 and lowered consumer prices across a meaningful chunk of the CPI basket. Growth for the fiscal year just closed came in strong, north of 7%, with the current fiscal year expected to moderate into the high sixes – still a very healthy print by global standards, just not quite the blistering pace of the year before.

A few other data points worth having in your head: the current account deficit has been narrowing, helped by strong services exports and remittance inflows that keep India as the largest recipient of private remittances globally, and gross FDI inflows have stayed robust, with Singapore, the US, Mauritius, the UAE, and the Netherlands accounting for the bulk of net inflows. None of that is spectacular, headline-grabbing news. It's exactly the kind of quiet, structurally sound macro backdrop that Dalio would want to see in the early-to-mid stage of a healthy internal cycle – sound finances, manageable inflation, a central bank with room to maneuver rather than one boxed into a corner.

What does a benign, low-inflation, steady-rate environment like this actually mean for this specific portfolio? For the lending book – Bajaj Finance, Cholamandalam, Muthoot, Aptus – a stable rate environment is close to ideal, because it means net interest margins aren't getting squeezed by a central bank forced into emergency hikes, and credit demand keeps growing into a economy still expanding at a healthy clip. For the consumption-adjacent names on the watchlist – DMart, Page Industries – a genuinely lower inflation print plus the GST rationalisation flowing through to consumer prices is a real tailwind for volume growth, because it protects real purchasing power at the household level. And for TCS and KPIT, the story is less about domestic rates and more about that resilient services-export and FDI picture – global capital continuing to see India as a preferred destination is the demand-side confirmation of the IT-sector thesis I laid out earlier.

I'll flag the honest risk too, because Marks would insist on it: the RBI's own commentary has pointed to persistently elevated energy prices and geopolitical tension as upside risks to inflation going forward, and the next policy review in early August will tell us whether the current calm holds or whether that risk starts showing up in the data. I'm not forecasting a shock. I'm just making sure I know what would have to happen for this benign backdrop to stop being benign.

The Boring Mechanics – Rebalancing, Cash, and Taxes

A few operational questions come up every quarter, and I'd rather answer them once, properly, than keep fielding them one at a time.

On rebalancing – I don't rebalance on a fixed calendar, and I think fixed-calendar rebalancing is actually a subtly bad habit, because it forces you to trim your winners and add to your losers purely on the basis of a date on a calendar, regardless of whether the underlying thesis for either position has actually changed. I'd rather let weights drift with genuine conviction and business performance, and only actively rebalance when a position's size has drifted for reasons that don't reflect improving fundamentals – pure multiple expansion untethered from the story, for instance. That's a judgment call, not a formula, and I know that's a less satisfying answer than a clean rule, but I'd rather give you an honest, judgment-based answer than a false sense of precision.

On cash – I run this portfolio close to fully invested most of the time, not because I think market timing is impossible, Taleb-style randomness arguments notwithstanding, but because I think the far more common mistake for someone my age and at this stage is holding excess cash out of a vague, unexamined anxiety rather than a specific, articulable reason. If I genuinely can't find eight names I have real conviction in, that's a different situation, and I'd rather hold cash and say so plainly than force a ninth mediocre position just to stay "fully invested" for its own sake. Right now, I have conviction in all eight, so cash stays minimal.

On taxes – this matters more than people give it credit for for a portfolio with genuinely low churn, because India's capital gains regime rewards patience explicitly: long-term gains on listed equity, held over a year, get taxed at a meaningfully lower rate than short-term gains. Every time I've been tempted toward a quick trade, running the actual after-tax math has been a useful, slightly humbling exercise, because a trade that looks attractive on a pre-tax basis often looks considerably less attractive once you account for both the tax drag and the transaction costs of moving in and out. Munger apparently used to say that a large fraction of his and Buffett's edge came simply from not paying capital gains tax repeatedly on the same pool of capital – compounding uninterrupted is a real, quantifiable advantage, not just a philosophical preference, and it's baked directly into why I'd rather hold a good-enough position through some discomfort than trade around it.

Why "Sovereign Anchor"

I don't think I've ever actually explained the name, and since some of you found this through the website, it seems worth doing properly.

"Sovereign" is the Dalio-and-Sharma half of the name – the recognition that no company, however well-run, escapes the

gravitational pull of the nation it operates in. Choose the soil first. That's not a slogan I picked because it sounded good; it's the literal first filter every single idea in this letter had to pass before I spent another hour on it.

"Anchor" is the other half, and it's meant almost literally – a portfolio that doesn't drift with every news cycle, every quarterly print, every macro headline. An anchor doesn't mean motionless; a ship at anchor still rises and falls with the tide, still gets pushed by the wind and the current. It just doesn't get swept away by any single gust. That's the eight-stock portfolio, I hope – genuinely responsive to real changes in the underlying businesses and the underlying country, structurally resistant to being yanked around by noise.

Put together, I think the name is honestly just a compressed version of the whole philosophy this letter has been trying to unpack for the last however-long-this-has-been: pick your country carefully, then hold your convictions steadily, and let time do the rest of the work.

Why Eight, and Not Five, or Fifteen

Someone asked me recently why eight specifically – why not round down to five, for even more concentration, or round up to fifteen, for a bit more cushion – and I think the honest answer is more interesting than the number itself.

Five would be a purer expression of conviction, and there's a real school of thought, one Munger himself apparently subscribed to at points in his career, that your best five ideas will always outperform your best twenty-five, simply because idea quality degrades as you're forced to go further down your own list. I take that argument seriously. But five also means a single company-specific accident – a fraud, a regulatory shock, a key-person risk turning out worse than anyone modeled – does proportionally more damage to the whole book than I'm currently comfortable with, given that I'm managing this as a genuine long-horizon compounding vehicle rather than a concentrated bet I can afford to blow up and rebuild.

Fifteen, on the other hand, starts to dilute the "minor business partner" framing I care about. I don't think I can genuinely track fifteen businesses at the depth I'd want to – reading every concall transcript, actually understanding the second-order effects of a management change or a regulatory shift – without the coverage becoming shallow theater rather than real partnership. That's Bevelin's point about the limits of individual judgment showing up in a very practical way: there's a ceiling on how many

things a single person can genuinely understand deeply, and pretending otherwise just means self-deceiving your way into shallow coverage disguised as diversification.

Eight is where those two pressures currently balance out for me – concentrated enough that each position has to earn real conviction, diversified enough that one bad surprise doesn't sink the whole book. I'll be honest that this isn't a number I derived from a formula. It's closer to a number I arrived at by noticing where my own attention started to genuinely thin out, and then giving myself a little room on top of that limit rather than running right up against it.

A Word on Risk Beyond the Balance Sheet

Everything I've told you about leverage, return on capital, and margin structure is balance-sheet risk, and I think it's the easiest kind of risk to talk about precisely because it's the most measurable. I want to spend a minute on the kind that doesn't show up on a balance sheet at all, because Munger's inversion habit – asking what could destroy this, rather than what could make it work – doesn't stop at the financial statements.

Regulatory risk is the first one, and it's genuinely live across a good chunk of this portfolio. NBFC risk-weight rules can change at the RBI's discretion in ways that affect Bajaj Finance, Chola mandalam, Muthoot, and Aptus simultaneously, regardless of how well any individual one of them is run. Gold-loan-specific regulation is a distinct risk sitting entirely inside Muthoot's business, given how central gold as collateral is to that model. None of this shows up as a line item anywhere. It shows up, if it shows up at all, as a sudden policy announcement that revalues an entire sub-sector overnight.

Currency and geopolitical risk is the second, and it's most concentrated in Caplin Point's US and Latin America exposure and in the broader IT sector's dependence on developed-market client budgets. A meaningful rupee move, or a shift in US healthcare or immigration policy that touches pharma supply chains, is the kind of tail risk that's genuinely hard to model with any precision – which is exactly why Taleb would tell me not to bother pretending I can model it precisely, and instead just make sure no single external shock of that kind can take down more than one position's worth of the portfolio at a time.

And the third is what I'd call key-person and succession risk, which is real and specific at the company level – Apollo's structure, for instance, still depends heavily on Dr. Reddy's now-adult daughters continuing to execute well across their

respective divisions, and any founder-led or founding-family business carries some version of this risk whether or not it's ever discussed openly in an annual report. I don't have a clean quantitative answer for any of these three categories. I just try to make sure I'm asking the question honestly each quarter, rather than only asking the questions that have tidy, spreadsheet-friendly answers.

The CFA and the Library – Two Different Kinds of Rigor

Since I mentioned being a CFA Level 3 candidate earlier, let me actually contrast that with everything I've just walked you through, because I think the difference is instructive, not just a resume detail.

The CFA curriculum is, at its best, an extremely thorough grounding in the standardized vocabulary and mechanics of finance – how to actually calculate a weighted average cost of capital, how derivatives are priced, how portfolio statistics behave, what the formal rules of ethics and professional conduct look like. It's rigorous in the way an engineering syllabus is rigorous. There's a right answer, and you either know the formula or you don't. I don't want to undersell that – it's given me a level of technical fluency I don't think I would have built on my own, reading books at my own pace, and when I'm building out a valuation model for something like GE Vernova or Apollo HealthCo, that formal training is doing real work underneath the narrative I give you.

But it's a genuinely different kind of rigor from what Taleb, or Dalio, or Bevelin are teaching, and I think conflating the two is a mistake a lot of freshly credentialed people make. The CFA curriculum, almost by necessity, teaches models that assume a cleaner, more statistically well-behaved world than the one we actually live in – it has to, because you can't test formal exam questions on genuinely irreducible uncertainty. The Library is where I go to be reminded, constantly, that the map is not the territory – that a discounted cash flow model is a useful structure for organizing your thinking, not a genuine window into the future, and that the moment you start treating your own model's precision as evidence of your own predictive accuracy, you've made exactly the mistake Taleb spent three books warning about.

So in practice, the two feed each other rather than compete. The CFA gives me the vocabulary and the mechanical toolkit to build a rigorous-looking case. The Library is what stops me from actually believing the case is more certain than it is, and reminds me to hold every model loosely enough that new evidence can still change

my mind. I'd trust neither discipline entirely on its own. Together, I think they're a genuinely useful combination – one teaches you how to build the model, the other teaches you how much weight the model actually deserves.

Looking Ahead to October

Let me close out the analytical portion of this by previewing what I'm watching for between now and the next dialogue, because I think a forward-looking checklist is more useful to you than another summary of what I've already told you.

On the portfolio itself: I want to see how Caplin Point's three Caplin 2.0 growth levers actually progress – the "Made in USA" margin expansion, the oncology ANDA pipeline, and Project 90's backward integration into API supply. Those three, taken together, are the difference between Caplin staying a steady compounder and genuinely re-rating upward, and I expect real data points on at least one of them before October.

On the watchlist: Maruti's SUV market-share trajectory under Vision 3.0 is the one I'm watching most closely, because it's the cleanest test of whether the multi-powertrain hedge is actually working commercially or just working on paper. I also want to see whether Apollo makes any concrete move toward separately listing the HealthCo digital arm, because that would meaningfully change how I think about valuing the whole group. And on the industrials side, I want to see whether GE Vernova T&D's order backlog conversion actually shows up in revenue growth that justifies the multiple the market's already paying for it, or whether that valuation gets tested.

On the macro side: the RBI's early-August policy meeting is the next real data point on whether the current benign, low-inflation backdrop holds, particularly given the central bank's own flagged concern about energy prices and geopolitical risk. And globally, I want to keep tracking whether the AI-productivity gap between frontier economies and diffusion economies like India starts showing up clearly in hard data – that's the one non-consensus call in this entire letter I'd revise fastest if the evidence started pointing the other way.

And on the research side, purely for my own development: I want to push further into evaluating genuinely foreign soil with the same rigor I'd apply to an Indian company, using Switzerland and Australia as the template rather than the destination. Not because I'm planning to deploy capital abroad soon, but because Sharma's whole point about reform cycles rising and falling is easier to spot in hindsight than in real time, and the only way to get

better at spotting it in real time is to keep practicing on countries where I don't yet have the answer memorized.

A Plain-Language Word on Risk

Before I get to the actual close, I want to say something about risk directly, in my own words, rather than tucking a disclaimer into fine print somewhere you'll never read it.

Everything in this dialogue is exactly that – a dialogue. It's how I'm thinking, and why, about my own capital. It is not a recommendation that you buy Caplin Point, or Bajaj Finance, or anything else I've mentioned, and it's especially not a recommendation on any of the roughly twenty watchlist names, most of which I don't own for reasons I've tried to spell out honestly along the way. My circumstances, my time horizon, my risk appetite, and my ability to sit through a genuinely bad quarter without flinching are mine, not yours, and Taleb's whole point about silent evidence applies here too – you're hearing about my process, not about the identical process run by someone whose eight stocks didn't work out, because that person isn't the one recording a podcast about it.

Markets go down as well as up, sometimes for reasons that have nothing to do with the quality of any individual business, and concentration – which I've spent a good chunk of this letter defending – cuts both ways by definition. Eight names means more upside if I'm right about all eight. It also means considerably more pain than a diversified index fund if I'm wrong about even two or three of them at the same time, particularly given how correlated the lending names are in a genuine stress scenario. I'd rather you hear that plainly from me than discover it only in a quarter where it's actually happening.

If You Want to Start Where I Did

A few of you have asked, after previous dialogues, where to actually begin if the Library section resonated and you wanted to build your own version of it rather than just borrowing mine secondhand. I won't pretend there's one correct order, but if I were doing it again from scratch, I'd start with Bevelin's *Seeking Wisdom*, because it's less a finance book than a manual for thinking clearly across any domain, and everything else on the shelf becomes easier to absorb once that foundation is in place. From there, I'd go to Taleb's *Fooled by Randomness* before *The Black Swan* or *Antifragile* – it's the gentlest entry point into the randomness argument, and the later two books build directly on it. Howard Marks's *The Most Important Thing* is short enough to read in

a weekend and dense enough to reread every year, which is exactly the combination you want in a book you're using as a working tool rather than a one-time read.

Once those three are genuinely absorbed, not just read, I'd branch toward whichever half of the shelf matches what you're actually trying to get better at — Dalio and Sharma if it's macro and country-level thinking you want to build, or Damodaran and Greenblatt and Mukherjea if it's company-level analysis and valuation you're after. I don't think you need all twenty-something books on my shelf to think well. I think you need three or four, read properly, argued with in the margins, and revisited when the market gives you a reason to test whether you actually understood them the first time or just agreed with them in the abstract.

Closing Thoughts

So that's the July dialogue — longer than anything I've given you before, and deliberately so, because I wanted this one to actually hold the full shape of how I've been thinking over the last three months, not a compressed highlight reel of it.

Let me try to pull the threads back together into something you can actually carry with you. The portfolio is eight names now, concentrated by design, weighted toward Indian financials because I think credit penetration in this country still has a long, unglamorous runway left, balanced against software and healthcare and engineering so the whole book isn't a single undiversified bet wearing four different tickers. The philosophy underneath it is borrowed, openly and gratefully, from people who thought about risk and uncertainty more carefully than I could have on my own — Taleb's humility about randomness, Marks's insistence on thinking past the first, obvious take, Dalio and Sharma's patient, decades-long lens on how nations actually rise and fall, Munger's habit of disqualifying an idea before falling in love with it, Damodaran's refusal to let a good story stand in without numbers to back it, and Mukherjea's reminder that the most unusual thing a business can do, in a market full of shortcuts, is simply compound honestly for a very long time.

This quarter's research took me to two countries that will never appear in this portfolio directly — Switzerland and Australia — because I think understanding what genuinely well-run soil looks like, in two very different forms, sharpens every judgment I make about the one country whose soil I've actually bet on. And it took me through roughly twenty companies and sectors I don't yet own, not as a preview of what's coming next, but as the ordinary, unglamorous discipline of staying in practice — reading widely,

rejecting often, so that the rare idea that actually deserves a ninth position doesn't slip past me because I wasn't paying attention.

I also tried to be honest about the parts of this that don't flatter me – the version of me who held a wider, less examined portfolio for reasons that had more to do with avoiding visible blame than managing genuine risk. I think that honesty matters more than I used to give it credit for, and if this merged format keeps that honesty intact going forward instead of letting it get edited out into two cleaner, more separate-sounding documents, then I think the experiment will have been worth it, regardless of how long it takes to record.

If you've made it this far – genuinely, thank you. Two hours is a real ask of anyone's time, and the fact that you gave it to hearing one 22-year-old think out loud about eight stocks and two countries he doesn't own is not something I take lightly. I'm not going to pretend I'll get every call in this letter right. Some of these watchlist names will get disqualified before they ever earn a real position. Some part of the current thesis on India's IT sector, or on Australia's productivity problem, or on Switzerland's cost base, will probably look naive in hindsight, the way most confident-sounding forecasts eventually do. That's fine. That's what Taleb would expect, and honestly, what I'd expect from myself too, three years into this.

What I can actually promise you is the process – reading widely, rejecting aggressively, sizing positions with real thought rather than habit, staying humble about what a 22-year-old with three years of experience genuinely knows versus what he's borrowed from people who've spent decades on this. I'd rather you judge me on that process, quarter after quarter, than on any single call in this letter landing perfectly.

I'll see you in October, and I suspect that one will merge research and dialogue again too, now that I've actually tried it and – at least from where I'm sitting – liked how honest it forced me to be.

Take care of yourselves, and I'll talk to you soon.